

Cornell Bowers C-IS
College of Computing and Information Science

Deep Learning



Week 7: Diffusion Models

Varsha Kishore, Justin Lovelace, Andrew Chang, Adhitya Polavaram



VAE Recap

Auto-Encoder

Input Manifold:
(intrinsically low dimensional)



Latent space:
(low dimensional)





Reconstruction Loss:

$$\min_{\phi, \theta} \sum_{\mathbf{x} \in \mathcal{D}} (\mathbf{x} - d_{\theta}(e_{\phi}(\mathbf{x})))^2$$



Auto-Encoder

Change 1: Latent *distributions*

Input Manifold:
(intrinsically low dimensional)



Latent space:
(low dimensional)




Reconstruction Loss:

$$\min_{\phi, \theta} \sum_{\mathbf{x} \in \mathcal{D}} (\mathbf{x} - d_{\theta}(e_{\phi}(\mathbf{x})))^2$$


Auto-Encoder

Change 1: Latent *distributions*

Input Manifold:
(intrinsically low dimensional)



x




mean
 μ



s


log-variance

Latent space:
(low dimensional)



$$\epsilon \sim \mathcal{N}(0, 1)$$

$$z = \mu + \epsilon s$$




x'



Reconstruction Loss:

$$\min_{\phi, \theta} \mathbb{E}_{z \sim q_{\phi}(z|x)} [(\mathbf{x} - d_{\theta}(z))^2]$$


Auto-Encoder

Change 2: Latent *prior* $\mathbf{P}(\mathbf{z})$ over latent spaceInput Manifold:
(intrinsically low dimensional) x 

e
encoder
 $q_{\phi}(\mathbf{z}|\mathbf{x})$

mean
 μ  s

log-variance

Latent space:
(low dimensional)

$p(z) = \mathcal{N}(0, I)$
prior

Reconstruction

$\epsilon \sim \mathcal{N}(0, 1)$
 $z = \mu + \epsilon e^s$



d
decoder
 $p_{\theta}(\mathbf{x}'|\mathbf{z})$

 x' 

$$\min_{\phi, \theta} \underbrace{\mathbb{E}_{z \sim q_{\phi}(\mathbf{z}|\mathbf{x})} \left[(\mathbf{x} - d_{\theta}(\mathbf{z}))^2 \right]}_{\text{reconstruction term}} + \underbrace{D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}) \| p(\mathbf{z}))}_{\text{prior matching term}}$$


Evidence Lower Bound (ELBO)

First note that we can minimize reconstruction loss or maximize rec. likelihood:

$$\min_{\phi, \theta} \underbrace{\mathbb{E}_{z \sim q_{\phi}(\mathbf{z}|\mathbf{x})} \left[(\mathbf{x} - d_{\theta}(\mathbf{z}))^2 \right]}_{\text{reconstruction term}} + \underbrace{D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}) \| p(\mathbf{z}))}_{\text{prior matching term}} = \max_{\phi, \theta} \underbrace{\mathbb{E}_{z \sim q_{\phi}(\mathbf{z}|\mathbf{x})} [\log p_{\theta}(\mathbf{x}|\mathbf{z})]}_{\text{reconstruction term}} - \underbrace{D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}) \| p(\mathbf{z}))}_{\text{prior matching term}}$$

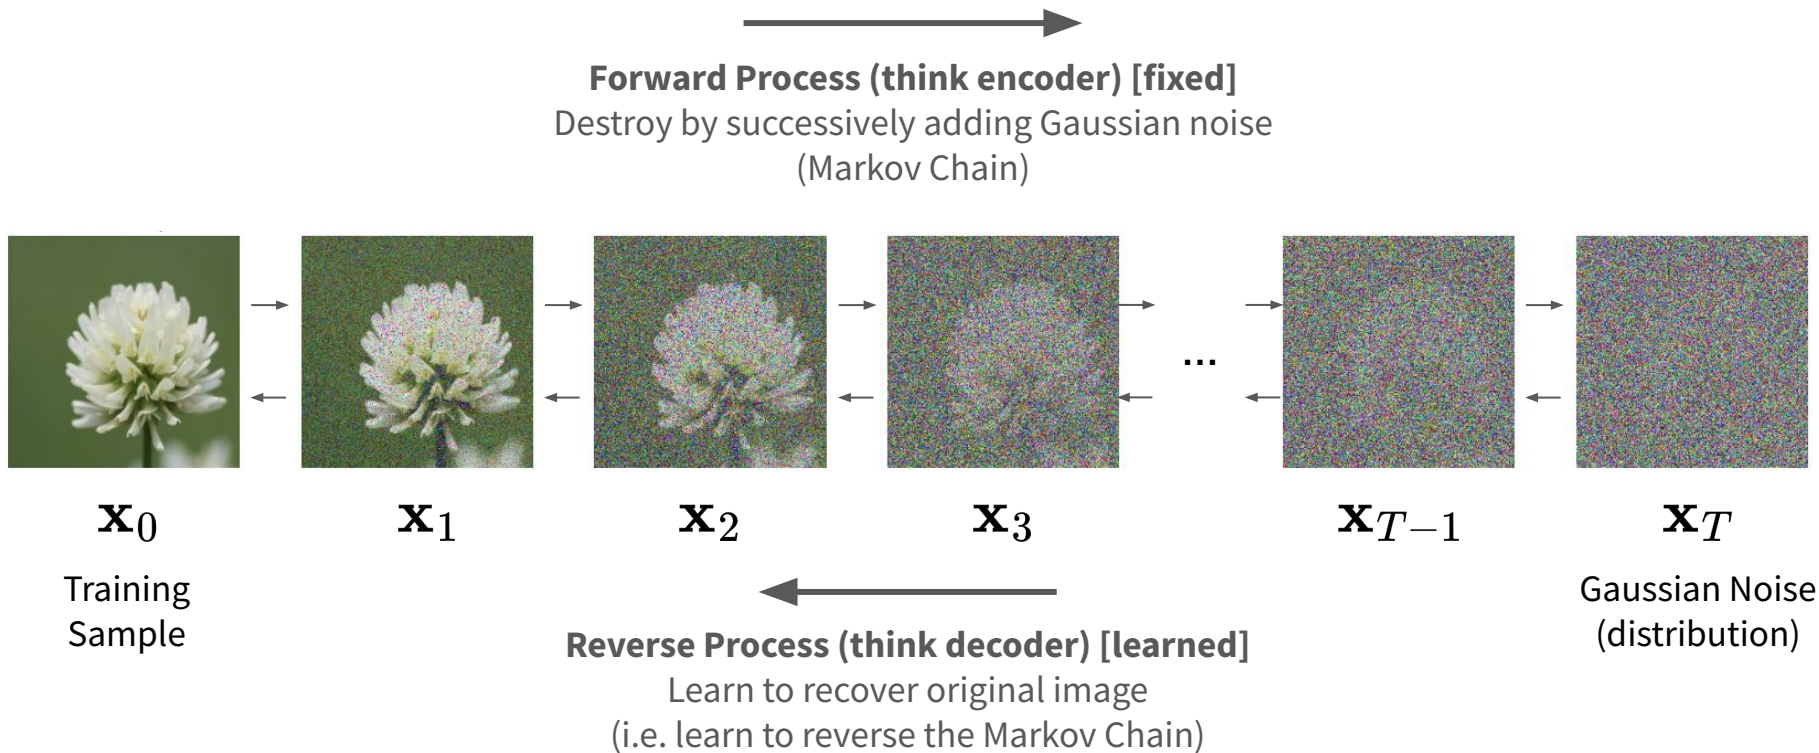
The ELBO proves that we are maximizing a lower bound of the data distribution:

$$\log(P(x)) \geq \underbrace{\mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [\log p_{\theta}(\mathbf{x}|\mathbf{z})]}_{\text{reconstruction term}} - \underbrace{D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}) \| p(\mathbf{z}))}_{\text{prior matching term}}$$



Diffusion Models

Denoising Diffusion Models





Forward Process

Markov Chain Implications

 $\mathbf{x}_0$  $\mathbf{x}_1$  $\mathbf{x}_2$  $\mathbf{x}_3$

...

 $\mathbf{x}_{T-1}$  $\mathbf{x}_T$

Direction of
dependence

Gaussian Noise: $q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{\alpha_t} \mathbf{x}_{t-1}, (1 - \alpha_t) \mathbf{I})$

$$q(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{x}_0) = q(\mathbf{x}_t | \mathbf{x}_{t-1}) \quad \text{T or F?}$$

$$q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) = q(\mathbf{x}_{t-1} | \mathbf{x}_t) \quad \text{T or F?}$$

What happens if
 $\alpha_{\square}=0$ and $\alpha_{\square}=1$?

Details: Forward Process


Can we extend this to sampling $\mathbf{x}_t$ in a closed form?

Let $\alpha_t := 1 - \beta_t$ $q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{\alpha_t} \mathbf{x}_{t-1}, (1 - \alpha_t) \mathbf{I})$

Re-parametrization trick!

$$\mathbf{x}_t = \sqrt{\alpha_t} \mathbf{x}_{t-1} + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_{t-1}$$

$$\boldsymbol{\epsilon}_{t-1} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$$




Brownian Motion towards the origin


$$\mathbf{x}_t = \sqrt{\alpha_t} \mathbf{x}_{t-1} + \sqrt{1 - \alpha_t} \epsilon_{t-1}$$

Interpret as two-step move:

1. Scale by $\times \sqrt{\alpha_t}$
2. Add noise $+\sqrt{1 - \alpha_t} \epsilon_t$



Details: Forward Process



Inductively, we can say

$$\mathbf{x}_t = \sqrt{\alpha_t} \mathbf{x}_{t-1} + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_{t-1}$$

$$= \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{1 - \alpha_t \alpha_{t-1}} \bar{\boldsymbol{\epsilon}}_{t-2}$$

Merged noise.
epsilon is still $\sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

$= \dots$

$$= \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}$$

$$\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$$




$$q(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\boxed{\phantom{\mathbf{x}_0}}, \boxed{\phantom{\mathbf{I}}})$$

We can sample $q(\mathbf{x}_t | \mathbf{x}_0)$ in closed form!!!

Reverse Process




The forward trajectory

$$q(\mathbf{x}_{0:T})$$



The reverse trajectory


$$p_{\theta}(\mathbf{x}_{0:T})$$



Reverse Process Overview

- "Learn to reverse what we just destroyed"
 - Learn time reversal of Markov Chain; we **train a model for this**
- Desired outcome: some $\mathbf{x}_0$ close to the original data distribution
- Problem: $q(\mathbf{x}_{t-1}|\mathbf{x}_t)$ is **intractable** (can't compute easily)

$$q(\mathbf{x}_{t-1}|\mathbf{x}_t) = \frac{q(\mathbf{x}_t|\mathbf{x}_{t-1})q(\mathbf{x}_{t-1})}{q(\mathbf{x}_t)}$$





Details: Reverse Process

$q(\mathbf{x}_{t-1}|\mathbf{x}_t)$ Is **not tractable**. Is $q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$ tractable?

I.e. Can we reverse the forward process given the original data?

$$q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) = \frac{q(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{x}_0)q(\mathbf{x}_{t-1}|\mathbf{x}_0)}{q(\mathbf{x}_t|\mathbf{x}_0)}$$

 $\mathbf{x}_T$  $\mathbf{x}_{T-1}$  $\mathbf{x}_2$  $\mathbf{x}_1$  $\mathbf{x}_0$

Key Idea

We introduce a generative model to approximate the reverse process $q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0)$



$$p(\mathbf{x}_T) = \mathcal{N}(\mathbf{0}, \mathbf{I})$$

$$p_{\theta}(\mathbf{x}_{t-1} | \mathbf{x}_t) = \mathcal{N}(\mu_{\theta}(\mathbf{x}_t, t), \sigma_t^2 \mathbf{I})$$

$$\mathbb{E}_{q(\mathbf{x}_t | \mathbf{x}_0)} [D_{\text{KL}}(q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) \parallel p_{\theta}(\mathbf{x}_{t-1} | \mathbf{x}_t))]$$

Learning Objective!

U-Net



$\mathbf{x}_T$

$\mathbf{x}_{T-1}$

$\mathbf{x}_2$

$\mathbf{x}_1$



$\mathbf{x}_0$

Reverse process:




The U-Net

[Ronneberger et al. 2015]




The U-Net

[Ronneberger et al. 2015]




Diffusion Sampling




$$p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t) = q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_\theta(\mathbf{x}_t, t)) \approx q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$$

Diffusion Sampling



Diffusion Sampling



Diffusion Models



An astronaut riding a horse

Algorithm 1 DDPM Training

- 1: **repeat**
- 2: Sample $\mathbf{x}_0 \sim q(\mathbf{x}_0)$
- 3: Sample $t \sim \mathcal{U}\{1, 2, \dots, T\}$
- 4: Sample $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 5: Compute noisy image:

$$\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$$

- 6: Take optimizer step on:

$$\mathcal{L}(\theta) = \mathbb{E}_{t, \mathbf{x}_0, \epsilon} [\|\epsilon - \epsilon_\theta(\mathbf{x}_t, t)\|^2]$$

- 7: **until** convergence
-

Algorithm 1 Deterministic Sampling (DDIM)

1: **Input:** trained noise predictor ϵ_θ , noise schedule $\alpha_1, \dots, \alpha_T$

2: Sample $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

3: **for** $t = T, T - 1, \dots, 1$ **do**

4: Predict noise:

$$\hat{\epsilon} = \epsilon_\theta(\mathbf{x}_t, t)$$

5: Estimate $\mathbf{x}_0$:

$$\hat{\mathbf{x}}_0 = \frac{\mathbf{x}_t - \sqrt{1 - \bar{\alpha}_t} \hat{\epsilon}}{\sqrt{\bar{\alpha}_t}}$$

6: Compute $\mathbf{x}_{t-1}$ deterministically:

$$\mathbf{x}_{t-1} = \sqrt{\bar{\alpha}_{t-1}} \hat{\mathbf{x}}_0 + \sqrt{1 - \bar{\alpha}_{t-1}} \hat{\epsilon}$$

7: **end for**

8: **return** $\mathbf{x}_0$

Q: Why is predicting the noise better than predicting the image?




ELBO for Diffusion Models

$$\begin{aligned}
 \log p(\mathbf{x}) &\geq \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \right] \\
 &= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T)p_\theta(\mathbf{x}_0|\mathbf{x}_1)}{q(\mathbf{x}_1|\mathbf{x}_0)} + \log \prod_{t=2}^T \frac{p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{\frac{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)q(\mathbf{x}_t|\mathbf{x}_0)}{q(\mathbf{x}_{t-1}|\mathbf{x}_0)}} \right] \\
 &= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T)p_\theta(\mathbf{x}_0|\mathbf{x}_1)}{\cancel{q(\mathbf{x}_1|\mathbf{x}_0)}} + \log \frac{\cancel{q(\mathbf{x}_1|\mathbf{x}_0)}}{q(\mathbf{x}_T|\mathbf{x}_0)} + \log \prod_{t=2}^T \frac{p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)} \right] \\
 &= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T)p_\theta(\mathbf{x}_0|\mathbf{x}_1)}{q(\mathbf{x}_T|\mathbf{x}_0)} + \sum_{t=2}^T \log \frac{p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)} \right] \\
 &= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} [\log p_\theta(\mathbf{x}_0|\mathbf{x}_1)] + \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T)}{q(\mathbf{x}_T|\mathbf{x}_0)} \right] + \sum_{t=2}^T \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)} \right] \\
 &= \mathbb{E}_{q(\mathbf{x}_1|\mathbf{x}_0)} [\log p_\theta(\mathbf{x}_0|\mathbf{x}_1)] + \mathbb{E}_{q(\mathbf{x}_T|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T)}{q(\mathbf{x}_T|\mathbf{x}_0)} \right] + \sum_{t=2}^T \mathbb{E}_{q(\mathbf{x}_t, \mathbf{x}_{t-1}|\mathbf{x}_0)} \left[\log \frac{p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)} \right] \\
 &= \underbrace{\mathbb{E}_{q(\mathbf{x}_1|\mathbf{x}_0)} [\log p_\theta(\mathbf{x}_0|\mathbf{x}_1)]}_{\text{reconstruction term}} - \underbrace{D_{\text{KL}}(q(\mathbf{x}_T|\mathbf{x}_0) \parallel p(\mathbf{x}_T))}_{\text{prior matching term}} - \sum_{t=2}^T \underbrace{\mathbb{E}_{q(\mathbf{x}_t|\mathbf{x}_0)} [D_{\text{KL}}(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \parallel p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t))]}_{\text{denoising matching term}}
 \end{aligned}$$

Training Objective


- Bound the likelihood with the ELBO
 - Exactly like VAEs

$$\log p(\mathbf{x}) \geq \underbrace{\mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})} [\log p_\theta(\mathbf{x}|\mathbf{z})]}_{\text{reconstruction term}} - \underbrace{D_{\text{KL}}(q_\phi(\mathbf{z}|\mathbf{x}) \parallel p(\mathbf{z}))}_{\text{prior matching term}}$$




$$\log p(\mathbf{x}) \geq \underbrace{\mathbb{E}_{q(x_1|x_0)} [\log p_\theta(x_0|x_1)]}_{\text{reconstruction term}} - \underbrace{D_{\text{KL}}(q(x_T|x_0) \parallel p(x_T))}_{\text{prior matching term}} - \sum_{t=2}^T \underbrace{\mathbb{E}_{q(x_t|x_0)} [D_{\text{KL}}(q(x_{t-1}|x_t, x_0) \parallel p_\theta(x_{t-1}|x_t))]}_{\text{denoising matching term}}$$

Diffusion Models vs. VAEs vs. GAN



Latent Diffusion

Latent Diffusion Models offer Excellent Trade-off between Performance and Compute Demands



Combines:
Auto-Encoders
Discriminators
Diffusion


- ➔ LDM with appropriate regularization, compression, downsampling ratio and strong autoencoder reconstruction:
- **Computationally efficient** diffusion model in latent space (compression & lower resolution).
 - Yet **very high-performance** (latent diffusion + autoencoder + discriminator = ❤️).
 - **Highly flexible** (can adjust autoencoder for different tasks and data).

Classifier-Free Guidance

Problem: Diffusion models often ignore the label (i.e. for example the text prompt)

Idea: Give the direction of the conditioning information extra emphasis.

Significantly improves quality of conditional models. Used by practically every conditional diffusion model




Increasing
Guidance



Video Generation




Prompt: 3D animation of a small, round, fluffy creature with big, expressive eyes explores a vibrant, enchanted forest.



Prompt: A cat waking up its sleeping owner demanding breakfast.

Quiz: How did Prof. Andrew Owens do this?

https://dangeng.github.io/visual_anagrams/



Overview

- Recap
- Diffusion model overview
- Forward
- Reverse
- Training Objective