



Cornell Bowers CIS

College of Computing and Information Science

CS 4782: Introduction to Deep Learning

Course Staff



Prof. Wei-Chiu Ma
Instructor



Prof. Kilian Q. Weinberger
Instructor



Žiga Kovačič
head-TA



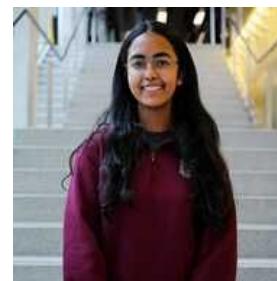
James Kim
Grad TA



Guangzhao He
Grad TA



Vibhhu Sharma
Grad TA



Akanksha Sarkar
Grad TA

Course Staff



Michael Wu
TA



Andrew Richmond
TA



Mac Turner
TA



Mohammad Labadi
TA



Rayhan Khanna
TA



William Leung
TA



Ram Peddu
TA



Sabrina Liu
TA



Chris Fernandes
Consultant



Daniel Lin
TA



Jane Lackley
TA



Haohuan Li
TA



Brandon Li
TA



Daniel Chuang
TA



Ethan Lin
TA



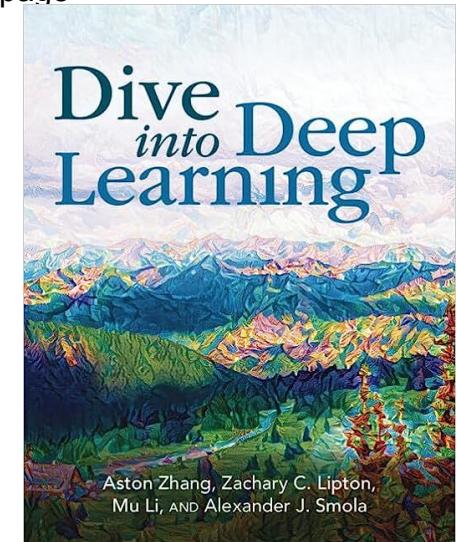
Jiasheng Shi
TA

Logistics

- All lectures will be held in person
- Lectures will be on Tuesdays and Thursdays from 2:55 to 4:10pm
- Please attend and participate!!

Logistics

- Course website: <https://www.cs.cornell.edu/courses/cs4782/>
 - Tentative **schedule**, homework policies, grading policies, etc. are on the course page
- Slides / **Office hours** are on the course website!
- No **course book**, but we will link to **DiDL** chapters
- Hub that links to everything: **Canvas page**
- Questions / Answers: **Ed Discussion**
- Projects: **Google Colab** (You will get free credits)
- Notes will usually be printed



Rules

- Never email the instructors directly
 - Post privately on **Ed Discussion**
 - Or email CS4782SP26@gmail.com
- No **laptops/mobiles/smart devices** in class
- **Class Code of Conduct** applies
- **Final Projects:** teams up to 5
- **Homework / Programming Projects:** teams up to 2



The screenshot shows a web browser window for 'Colab Paid Services Pricing' at colab.research.google.com/signup. The page title is 'Choose the Colab plan that's right for you'. It describes Colab as always free but offers paid options for increased computing needs. A link to 'Restrictions apply, learn more here' is provided.

Pay As You Go

- \$9.99 for 100 Compute Units
- \$49.99 for 500 Compute Units

You currently have 0 compute units. Compute units expire after 90 days. Purchase more as you need them.

- No subscription required. Only pay for what you use.
- Faster GPUs. Upgrade to more powerful GPUs.

Recommended

Colab Pro

\$9.99 per month

No cost for students and educators

- 100 compute units per month. Compute units expire after 90 days. Purchase more as you need them.
- Faster GPUs. Upgrade to more powerful GPUs.
- More memory. Access our highest memory machines.

Colab Pro+

\$49.99 per month

All of the benefits of Pro, plus:

- An additional 500 compute units per month, totaling **600** compute units per month.
- Compute units expire after 90 days. Purchase more as you need them.
- Faster GPUs. Priority access to upgrade to more powerful premium GPUs.
- Background execution. With compute units, your actively running notebook will continue running for up to 24hrs, even if you close your browser.

Colab Enterprise

Pay for what you use

- Integrated. Tightly integrated with Google Cloud services like BigQuery and Vertex AI.
- Enterprise notebook storage. Replace your usage of Google Drive notebooks with GCP notebooks, stored and shared within your cloud console.
- Productive. Generative AI powered code completion and generation.

Colab is offered by Google. For information on bulk pricing, please contact us.

Contact us

Grading (4782)

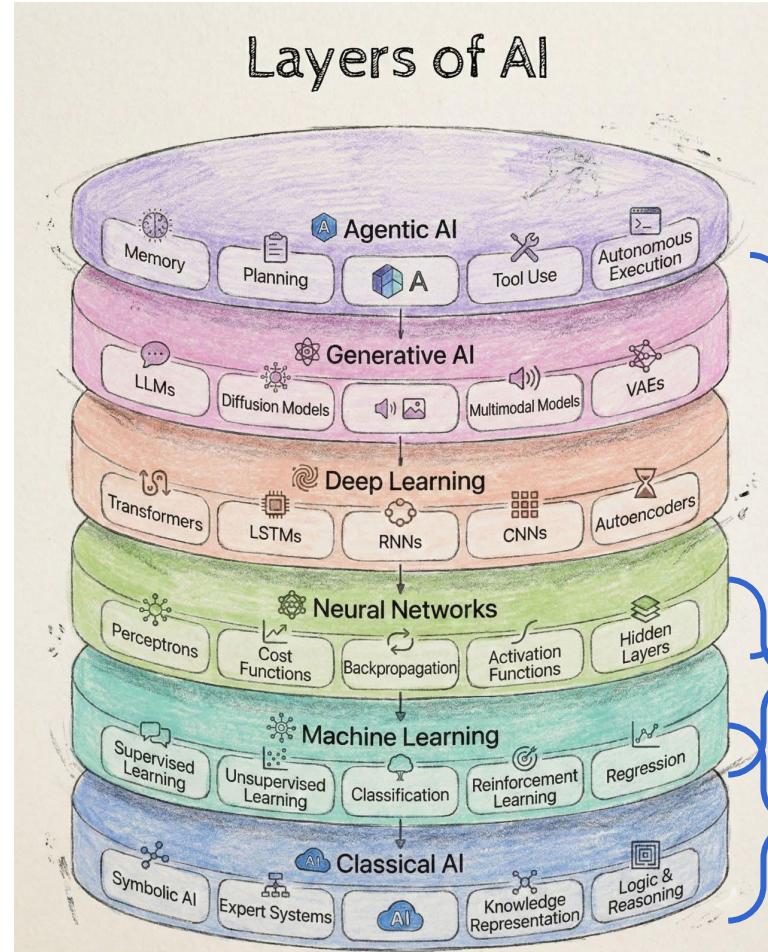
- Homework (10%)
 - There will be written assignments and coding projects
 - Google Cloud Credits for compute!
 - We recommend doing them in pairs!
 - 2-slip days for every assignment
 - **NEW in 2026: Only graded for completion!**
- Mid-term exam (40%)
 - Will be similar to the homework assignments
- Project (35%)
 - Goal: familiarize yourself with deep learning libraries
 - Implement a method from a recent research paper and reproduce their results
- Participation (15%)
 - Attend classes!
 - **NEW in 2026: Pop Quizzes**
 - Engage in class discussions and/or post on EdStem
 - Provide feedback to improve the class (we will reach out to you)

Grading (5782) [voluntary opt-in for 4782 students]

- Homework (10%)
 - There will be written assignments and coding projects
 - Google Cloud Credits for compute!
 - We recommend doing them in pairs!
 - 2-slip days for every assignment
 - **NEW in 2026: Only graded for completion!**
- Mid-term exam (35%)
 - Will be similar to the homework assignments
- Project (30%)
 - Goal: familiarize yourself with deep learning libraries
 - Implement a method from a recent research paper and reproduce their results
- Participation (15%)
 - Attend classes!
 - **New in 2026: Pop Quizzes**
 - Engage in class discussions and/or post on EdStem
 - Provide feedback to improve the class (we will reach out to you)
- **Paper Quizzes (10%)**
 - **Read specified research papers**
 - **Answer quizzes on Canvas**

Academic Integrity

- Do not disclose exact solutions to members from other groups for assignments
 - High-level discussion is allowed
- Cite any external sources
- You can use ChatGPT/Gemini/other AI assistants
 - But you **must add a note** explaining what you used it for and how you used it



ML/AI Courses at Cornell (incomplete)

CS 3700: Foundations of AI Reasoning and Decision-Making

CS 3780: Introduction to Machine Learning (prereq of CS4782)

CS 4782: Introduction to Deep Learning (this class)

CS 4670: Introduction to Computer Vision

CS 4740: Natural Language Processing

(superset of CS4782)

CS 4789: Introduction to Reinforcement Learning

CS 4756: Robot Learning

CS 4775: Computational Genetics and Genomics

Course Objectives

By the end of the course you will be able to...

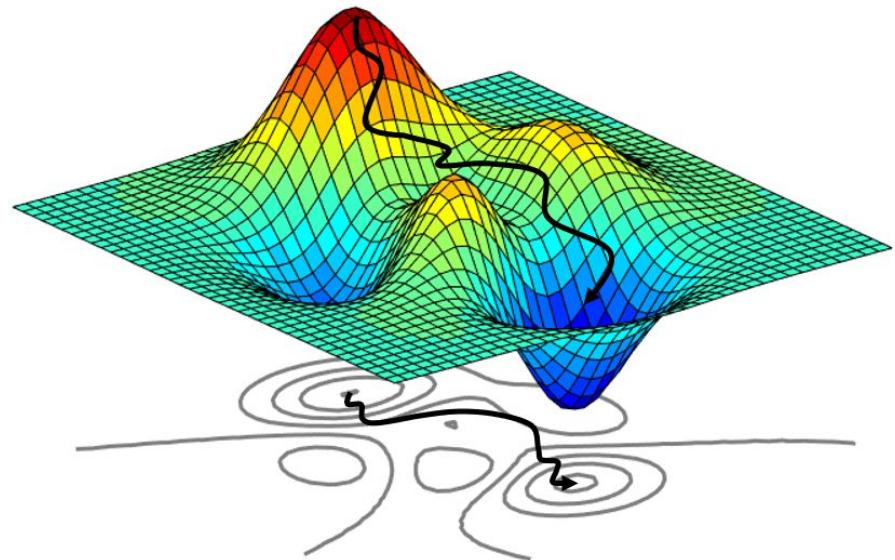
1. Design, train, and evaluate deep neural networks
2. Apply deep learning techniques to solve real-world problems in computer vision, natural language processing, and other complex domains
3. Critically evaluate pros/cons of different model architectures
4. Read and understand research in deep learning
5. Understand the core design principles behind leading deep learning systems like GPT-4, DALL-E 2/3, and Stable Diffusion

Topics Covered in CS4782



Neural Networks Review and Optimization

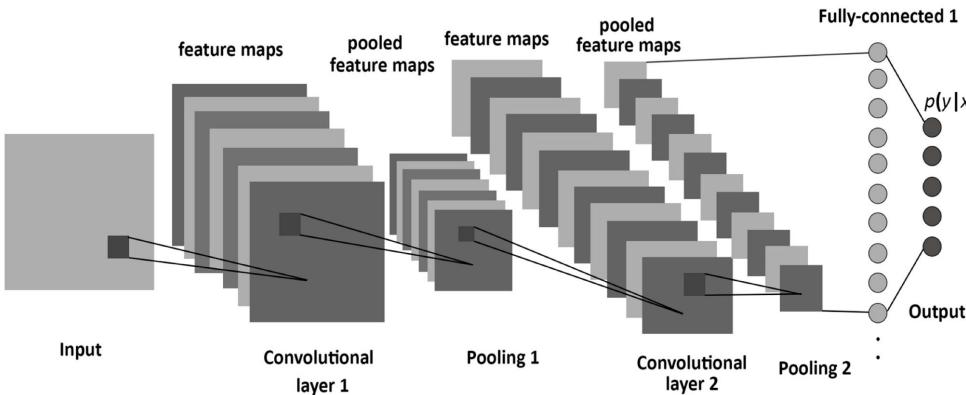
- Optimization algorithms - gradient descent, SGD, AdaGrad, Adam
- Learning rate scheduling
- Hyperparameter Optimization
- Regularization



<https://towardsdatascience.com/an-introduction-to-surrogate-optimization-intuition-illustration-case-study-and-the-code-5d9364aed51b>

Computer Vision

- Convolutional neural networks
- Different convolutional architectures - vanilla CNN, LeNet, ResNet, DenseNets



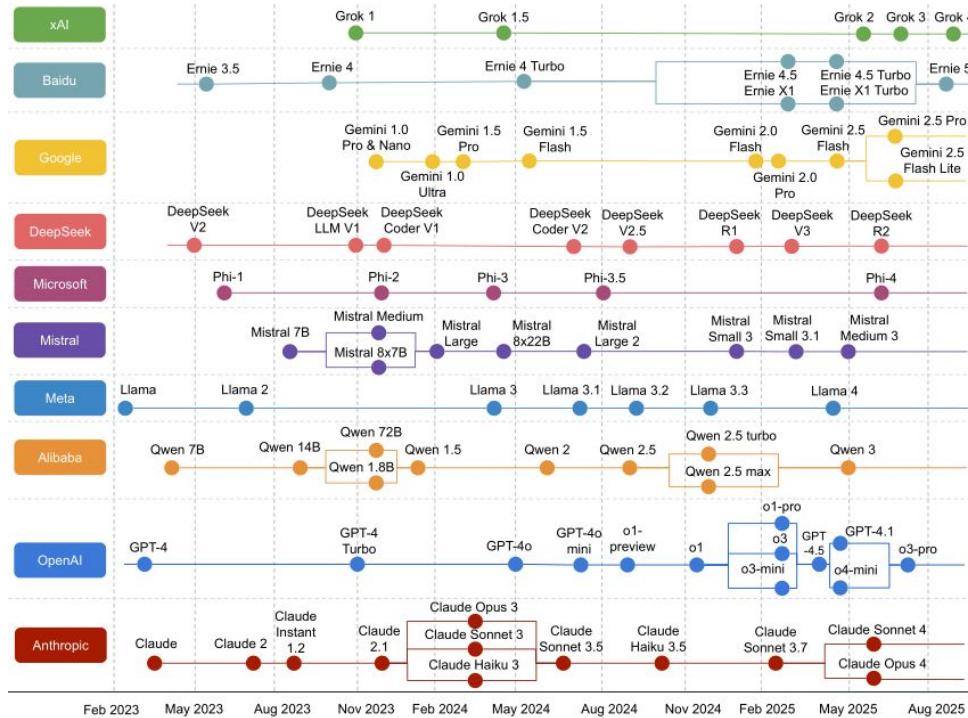
<https://www.mdpi.com/1099-4300/19/6/242>



CVPR 2018 WAD Video Segmentation

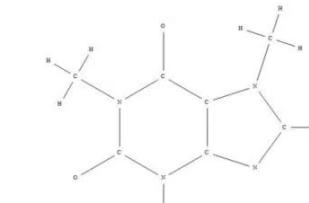
Natural Language Processing

- Word Embeddings
- Recurrent Neural Networks
 - RNNs/ LSTMs
- Attention and Transformers
- Large Language Models (LLMs)

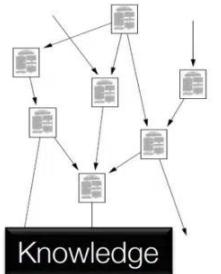


Graph Neural Networks

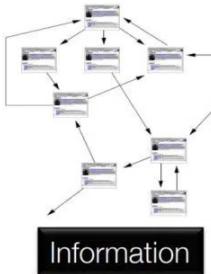
- Neural networks for data represented as graphs!



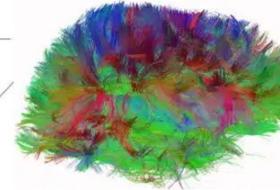
Molecules



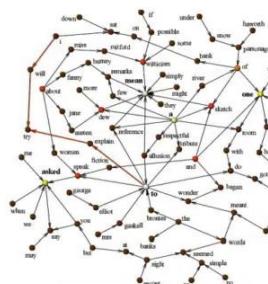
Knowledge



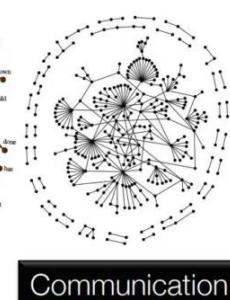
Information



Brain/neurons

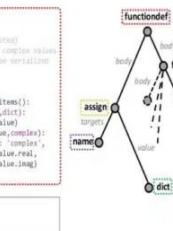


Genes



Communication

```
def encode(obj):  
    """Encode (possibly nested) dictionary containing complex values  
    into a form that can be serialized using JSON.  
  
    If the value is a list or dict:  
        - keys() & encode(value)  
    elif instance(value,complex):  
        e[key] = ('type': 'complex',  
                 'value':value.read,  
                 'l': value.lang)  
  
    return e  
  
import ast  
tree = ast.parse(  
    """  
    def functiondef():  
        body = [body]  
        for target in body:  
            if target == 'for':  
                body.append(target)  
            else:  
                target.append(body)  
        return body  
    """  
    )
```



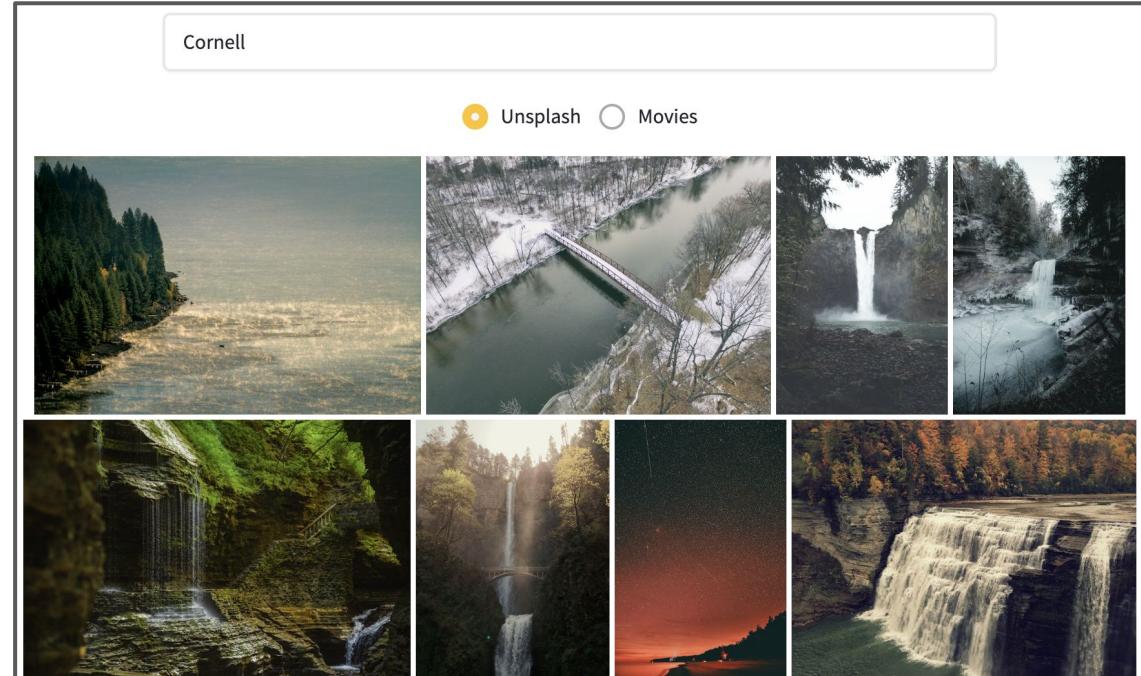
Software



Social

Modern Vision Networks

- Vision Transformers (ViTs)
- Vision Pre-Training
 - (Supervised, Self-supervised)
- Vision-Language Models

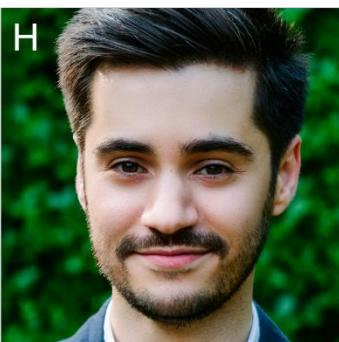
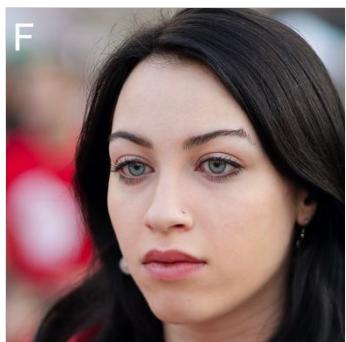
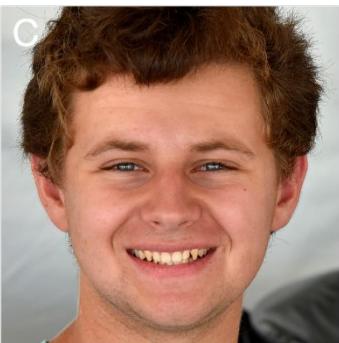
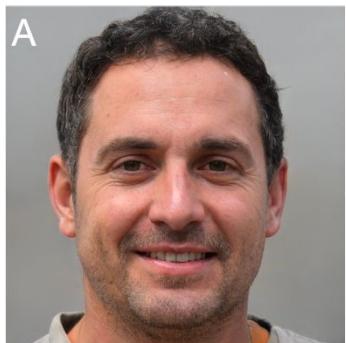


Generative Models

- U-Nets
- Variational Autoencoders (VAEs)
- Generative Adversarial Networks (GANs)
- Diffusion Models
- Multi-Modal Diffusion



Real or Fake?



Real or Fake?



Cornell Bowers CIS

2023 → 2026



GPT5.2: Make this image more realistic, the background and neck are weird.

Reinforcement Learning

Technique for an agent to learn in an interactive environment by testing different actions and obtaining feedback from its experiences.

- Markov Decision Process
- Q-learning/Deep Q-learning
- Policy Gradients
- Exploration strategies
- RL from Human Feedback

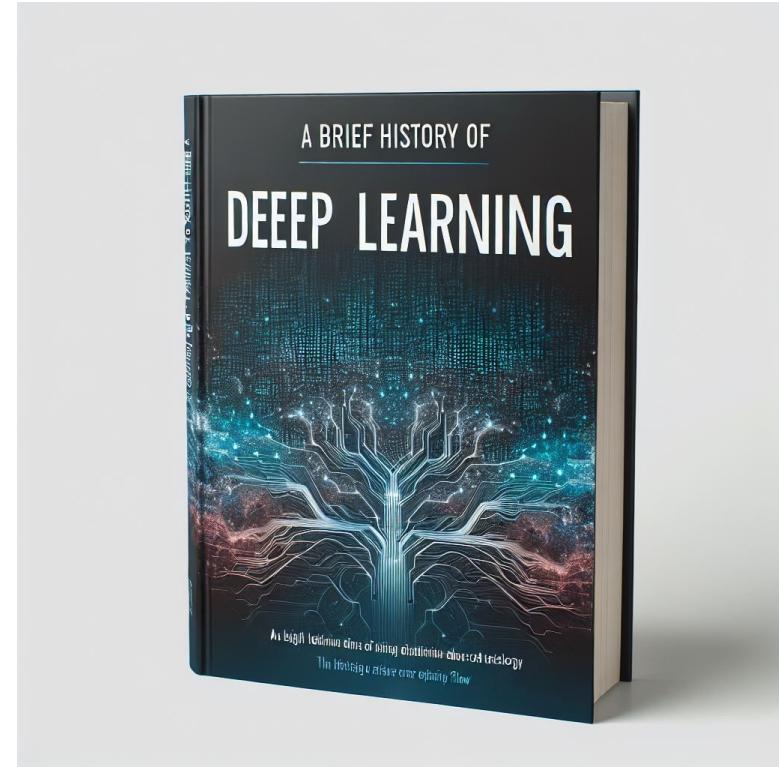


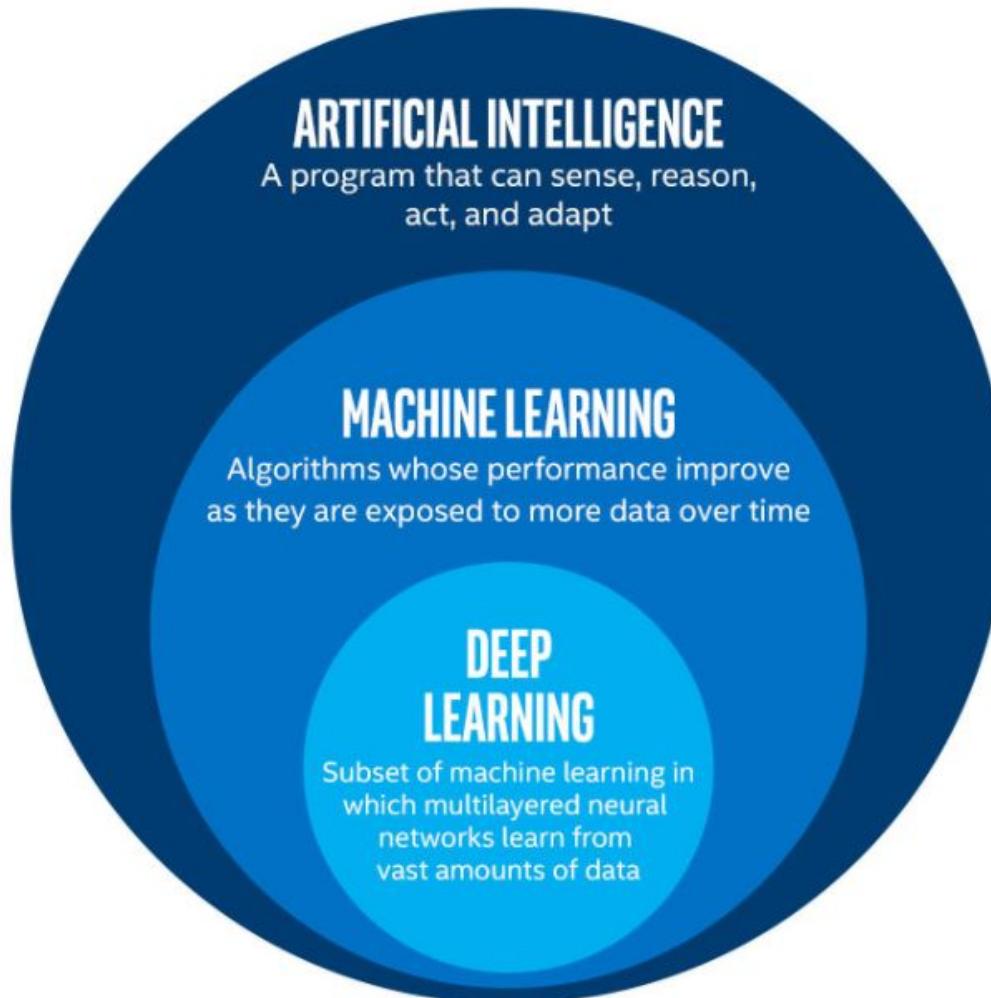
AI in Human Society

- Opportunities for mankind
- Ethical dilemmas
- Threats to humanity
 - Weapons
 - Loss of control
 - Mental Health
- Expected Winners & Losers
 - Job losses
 - Shift in Wealth
- Legislative Concerns
 - Accountability
 - Regulatory uncertainty
 - AI outpaces legislation



A brief history of Deep Learning





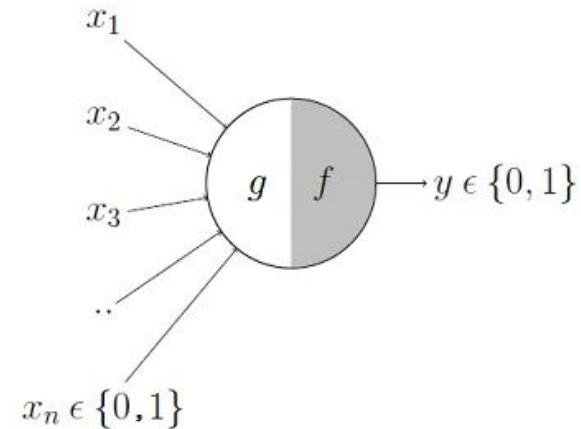
McCulloch-Pitts Neuron

Computational model of a neuron that was proposed by Warren McCulloch (neuroscientist) and Walter Pitts (logician) in 1943.

Showed that:

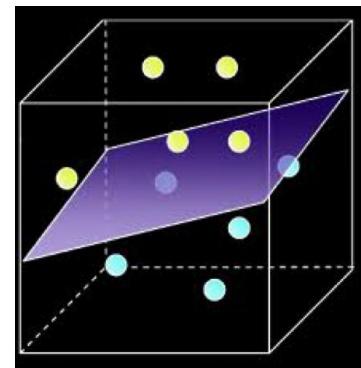
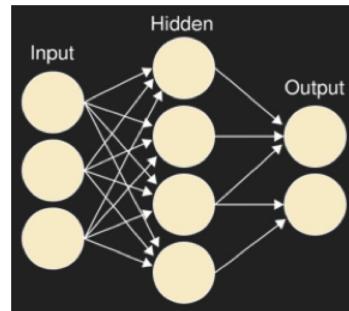
- Neurons can implement logical operations (AND, OR, NOT)
- Networks of neurons can compute any Boolean function
- The brain could be understood as a computational system

Purely mathematical model, did not **learn**, only allowed binary values

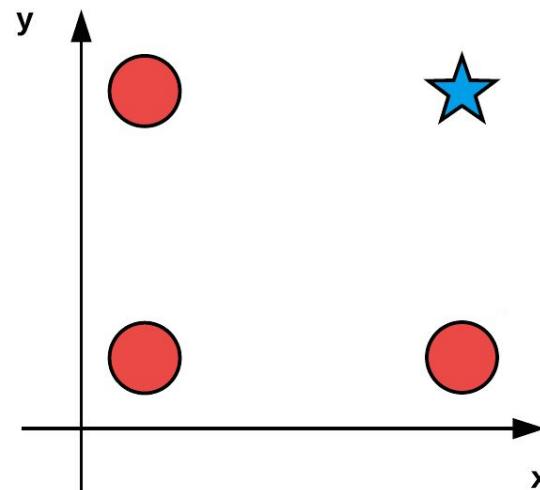


Perceptron (1957)

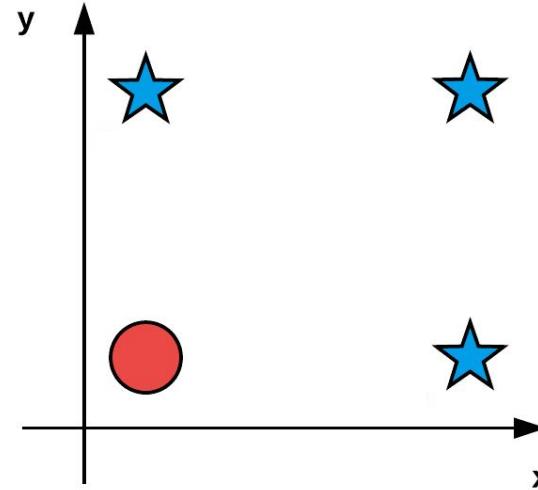
- Linear classifier, predecessor to Neural Network
- Trained with the perceptron update rule
- Invented @ Cornell University
 - First task: Recognize the Cornell “C” Logo
- Sparked **huge A.I.** excitement world wide



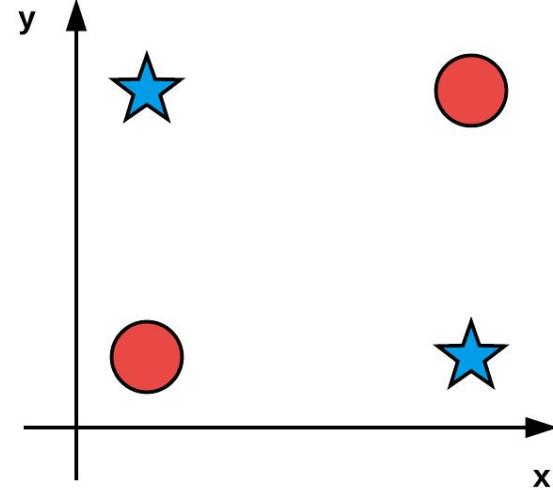
AND



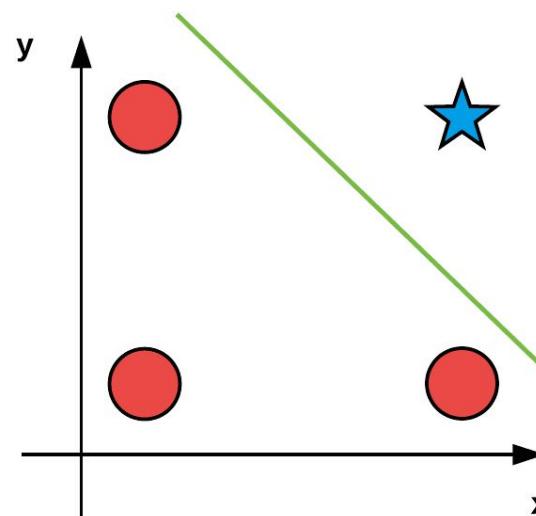
OR



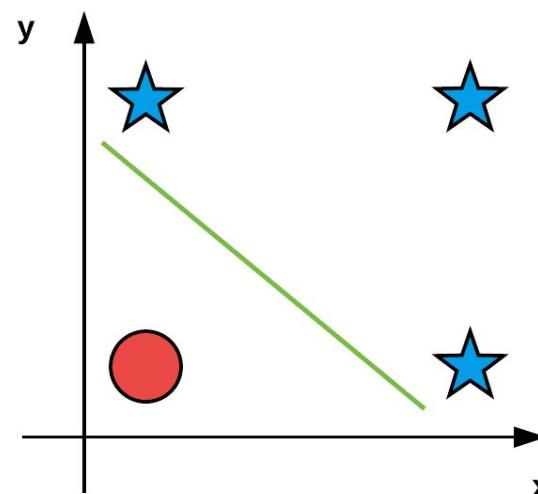
XOR



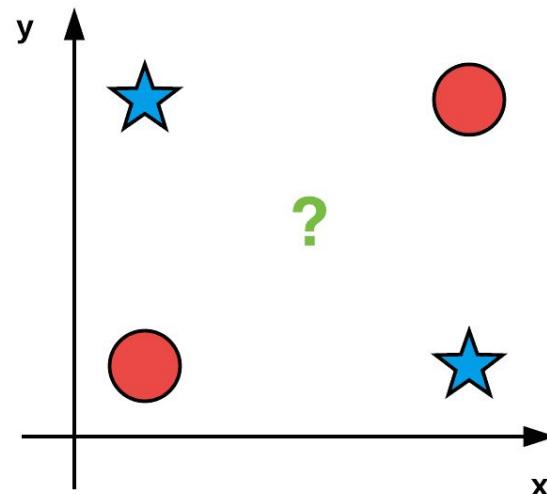
AND



OR

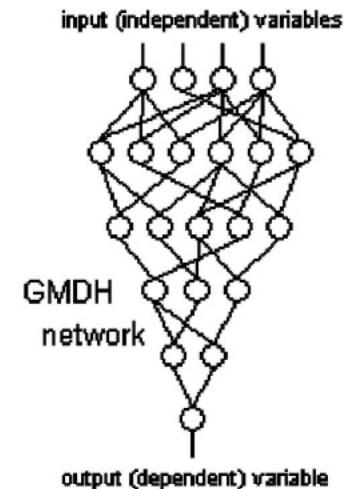


XOR



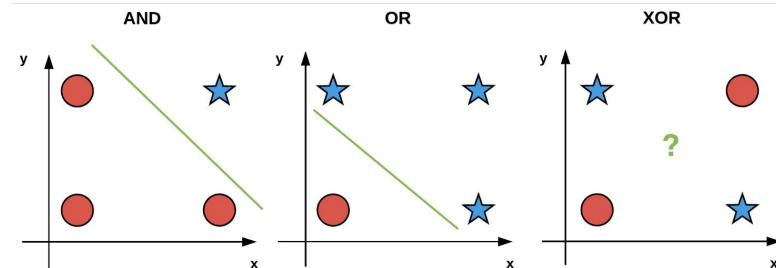
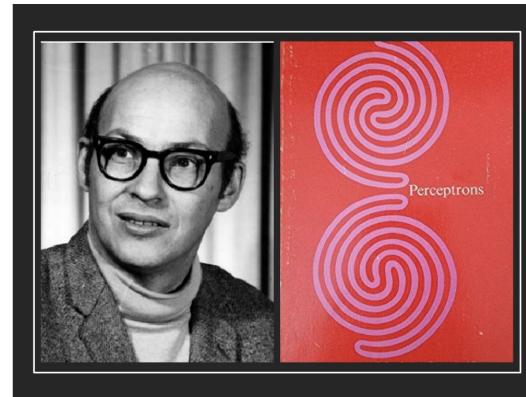
Multi-layer neural networks

- Multi-Layer Perceptron, Rosenblatt (around 1965)
- Alexey Grigoryevich Ivakhnenko 1965 Group Method of Data Handling (GMDH)
 - 1971 Eight Layer Neural Nets with skip connections!



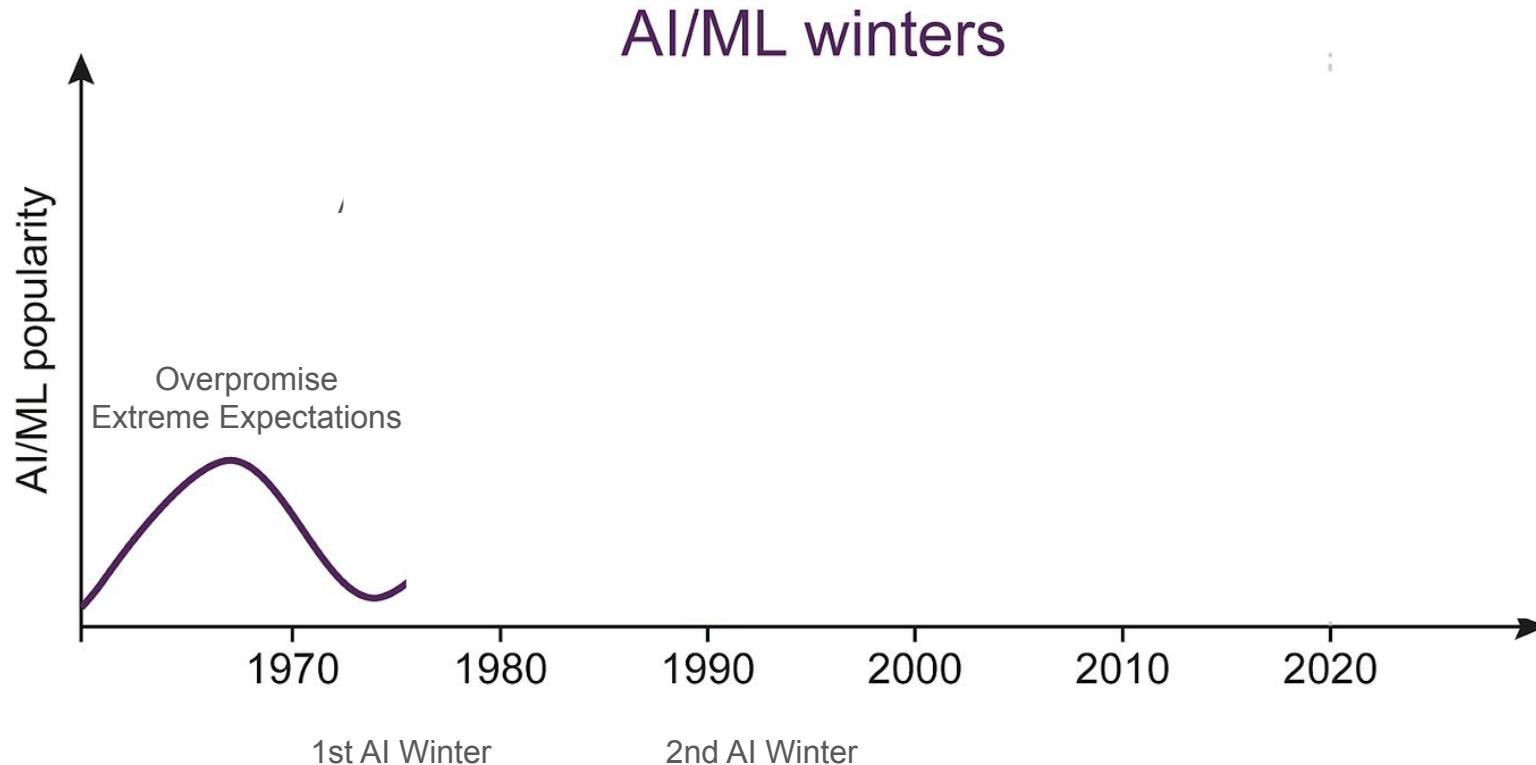
AI Winter (1974-1980)

- (1969) Minsky & Papert “killed” AI
- Burst huge expectation bubble
- Speech understanding / translation fails
- UK and US stop funding AI research



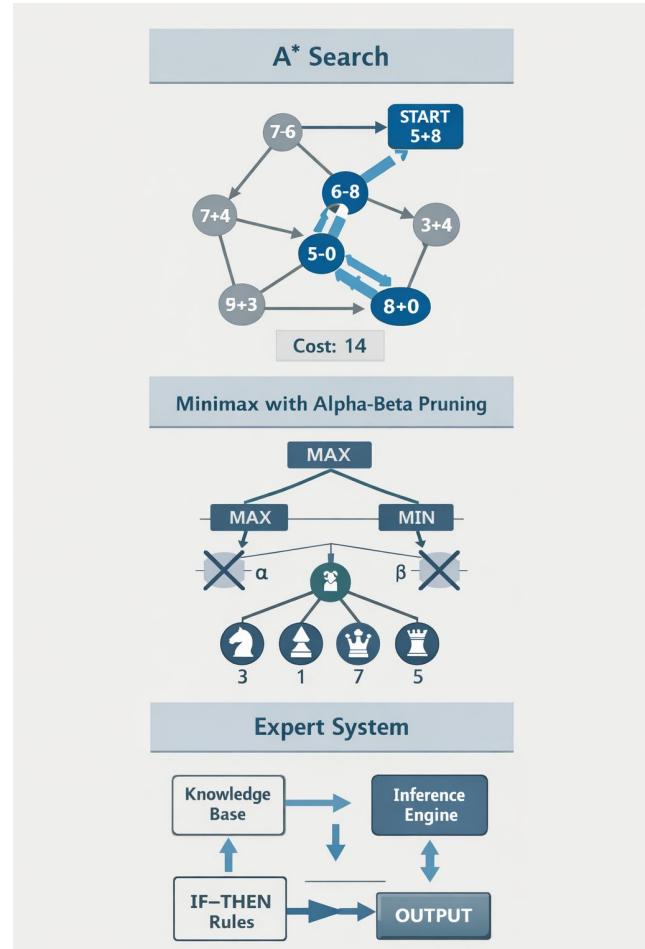
<https://www.pyimagesearch.com/2021/05/06/implementing-the-perceptron-neural-network-with-python/>

Public Perception of AI/ML



A.I. 2.0 - Search Algorithms

- **Search Algorithms**
 - A*: Optimal pathfinding using heuristics (cost + estimate)
 - Problem-solving framed as state-space search
- **Game Playing**
 - Minimax with alpha–beta pruning
 - Strong performance in chess and checkers
- **Expert Systems**
 - Rule-based systems: *IF–THEN* logic
 - Knowledge encoded from human experts
 - Widely used in medicine, engineering, and diagnostics

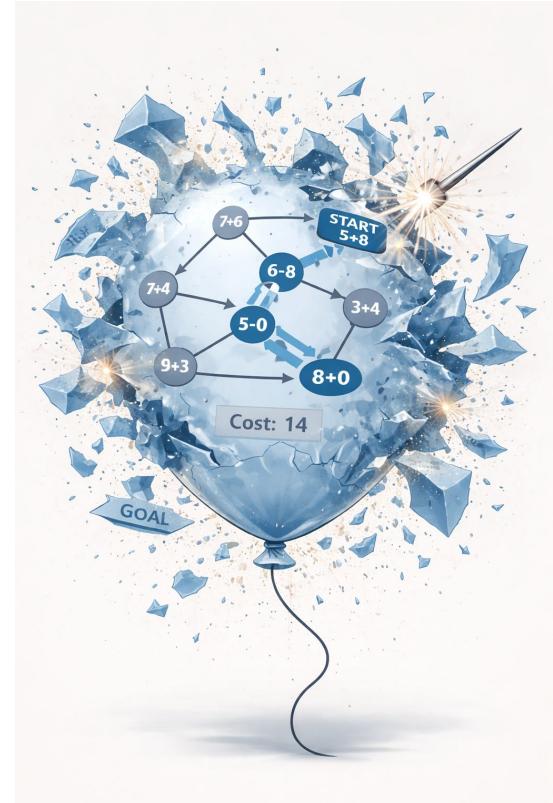


A.I. 2.0 Bubble Bursts!

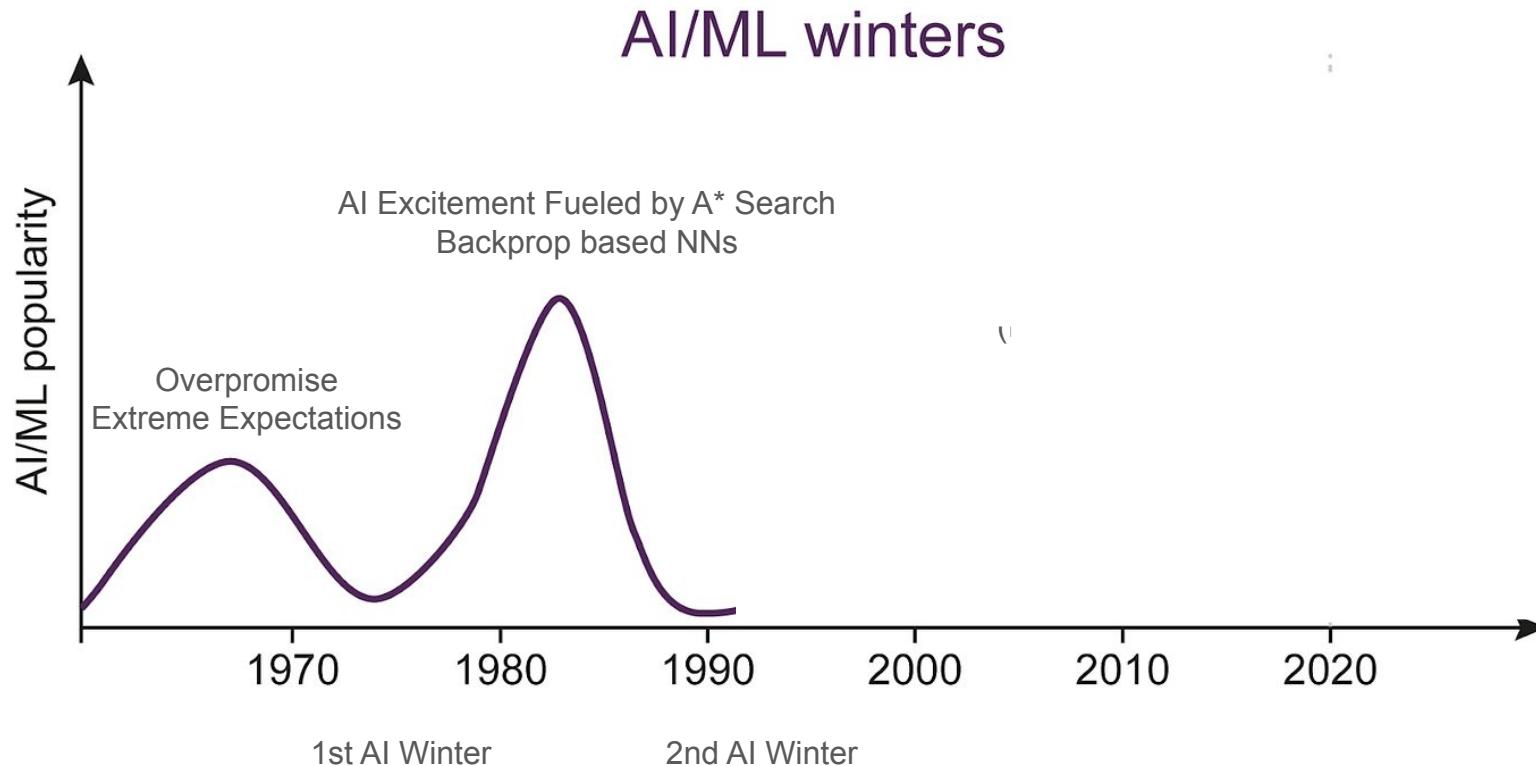
Limitations of search:

- Search trees scale **exponentially** with problem size
 - Toy problems (e.g. chess endgames) did not translate to real world
- Reliance on (brittle) hand-crafted heuristics
- Weak generalization across problems
 - No learning from data
 - No improvement with experience
- Poor handling of uncertainty / Noise
- Works within narrow assumptions
 - Minor deviations cause incorrect reasoning / nonsensical outputs

But search still has its place in modern A.I.

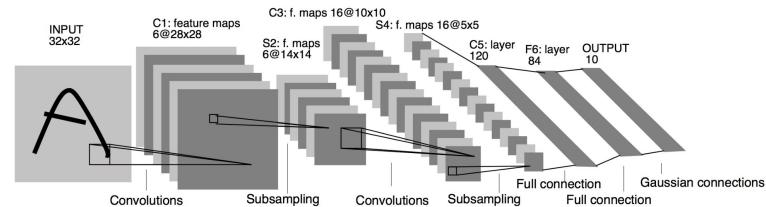
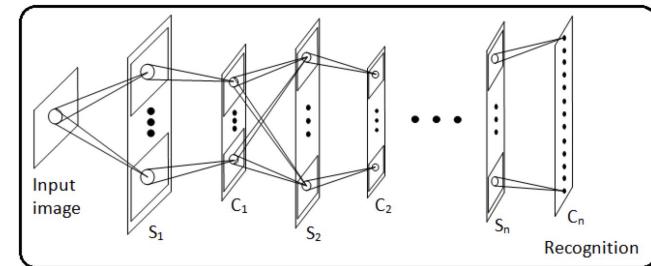


Public Perception of AI/ML



ConvNets

- 1979 Kunihiko Fukushima invents Neocognitron
 - Heavily inspired by human Visual Cortex
 - Alternates between Simple Cells / Complex Cells
 - Unsupervised
- 1986 Yann LeCun introduces BackProp to ConvNets for Handwritten Digits (creates MNIST)



Recurrent Neural Nets

- 1982 John Hopfield “Hopfield Networks”
- 1991 Sepp Hochreiter formulates Vanishing Gradient Problem
- 1997 S. Hochreiter and Jürgen Schmidhuber publish “Long Short-Term Memory” (LSTM)
 - <https://people.idsia.ch/~juergen/ai-priority-disputes.html>

Proc. Natl. Acad. Sci. USA
Vol. 79, pp. 2554–2558, April 1982
Biophysics

Neural networks and physical systems with emergent collective computational abilities

(associative memory/parallel processing/categorization/content-addressable memory/fail-safe devices)

J. J. HOPFIELD

Division of Chemistry and Biology, California Institute of Technology, Pasadena, California 91125, and Bell Laboratories, Murray Hill, New Jersey 07974

Contributed by John J. Hopfield, January 15, 1982

ABSTRACT Computational properties of use to biological organisms or to the construction of computers are described as collective properties of systems having a large number of roughly equivalent components (or neurons). The physical meaning of content-addressable memory is described by an appropriate phase space flow of the state of a system. A model of such a system is given, based on aspects of neurobiology but readily adapted to integrated circuits. The collective properties of this model produce a content-addressable memory which correctly yields an entire memory from any subpart of sufficient size. The algorithm for the time evolution of the state of the system is based on asynchronous parallel processing. Additional collective properties include tolerance for generalization, familiarity recognition, categorization, error correction, and time sequence retention. The collective properties are only weakly sensitive to details of the modeling or the failure of individual devices.

Given the dynamical electrochemical properties of neurons and their interconnections (synapses), we readily understand schemes that use a few neurons to obtain elementary useful biological behavior (1–3). Our understanding of such simple circuits in electronics allows us to plan larger and more complex circuits which are essential to large computers. Because evolution has no such plan, it becomes relevant to ask whether the ability of large collections of neurons to perform computational tasks may in part be a spontaneous collective consequence of having a large number of interacting simple neurons.

In physical systems made from a large number of simple elements, interactions among large numbers of elementary components yield collective phenomena such as the stable magnetic orientations and domains in a magnetic system or the vortex patterns in fluid flow. Do analogous collective phenomena in

calized content-addressable memory or categorizer using extensive asynchronous parallel processing.

The general content-addressable memory of a physical system

Suppose that an item stored in memory is “H. A. Kramers & G. H. Wannier *Phys. Rev.* **60**, 252 (1941).” A general content-addressable memory would be capable of retrieving this entire memory item on the basis of sufficient partial information. The input “& Wannier” (1941) might suffice. An ideal memory could deal with errors and retrieve this reference even from the input “Vanier” (1941). In contrast, a relatively simple form of content-addressable memory has been made in hardware (10, 11). Sophisticated ideas like error correction in accessing information are usually introduced as software (10).

There are classes of physical systems whose spontaneous behavior can be used as a form of general (and error-correcting) content-addressable memory. Consider the time evolution of a physical system that can be described by a set of general coordinates. A point in state space then represents the instantaneous condition of the system. This state space may be either continuous or discrete (as in the case of N Ising spins).

The equations of motion of the system describe a flow in state space. Various classes of flow patterns are possible, but the system's use for memory particularly include those that flow toward locally stable points from anywhere within regions around those points. A particle with frictional damping moving in a potential well with two minima will exhibit such a dynamics.

If the flow is not completely deterministic, the description is more complicated. In the two-well problems above, if the frictional force is characterized by a temperature, it must also produce a random driving force. The limit points become small

Universal Approximation

- 1989 George Cybenko proofs universal approximation of single hidden-layer neural networks
- Also yields wide-spread believe that more than one layer is unnecessary

Math. Control Signals Systems (1989) 2: 303–314

Mathematics of Control,
Signals, and Systems
© 1989 Springer-Verlag New York Inc.

Approximation by Superpositions of a Sigmoidal Function*

G. Cybenko†

Abstract. In this paper we demonstrate that finite linear combinations of compositions of a fixed, univariate function and a set of affine functionals can uniformly approximate any continuous function of n real variables with support in the unit hypercube, only mild conditions are imposed on the univariate function. Our results settle an open question about representability in the class of single hidden layer neural networks. In particular, we show that arbitrary decision regions can be arbitrarily well approximated by continuous feedforward neural networks with only a single internal, hidden layer and any continuous sigmoidal nonlinearity. The paper discusses approximation properties of other possible types of nonlinearities that might be implemented by artificial neural networks.

Key words. Neural networks, Approximation, Completeness.

1. Introduction

A number of diverse application areas are concerned with the representation of general functions of an n -dimensional real variable, $x \in \mathbb{R}^n$, by finite linear combinations of the form

$$\sum_{j=1}^N \alpha_j \sigma(y_j^T x + \theta_j), \quad (1)$$

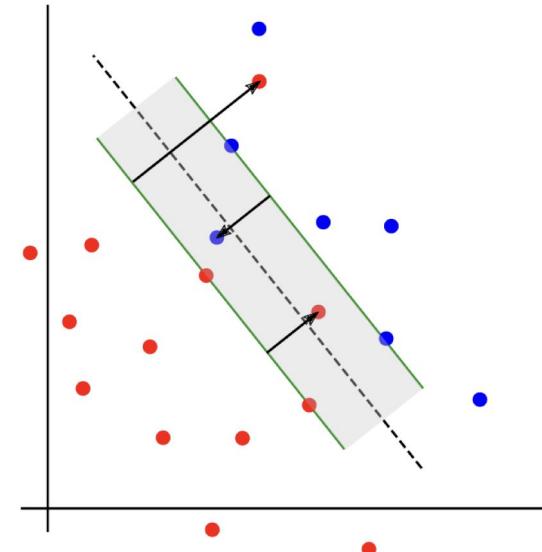
where $y_j \in \mathbb{R}^n$ and $\alpha_j, \theta \in \mathbb{R}$ are fixed. (y^T is the transpose of y so that $y^T x$ is the inner product of y and x .) Here the univariate function σ depends heavily on the context of the application. Our major concern is with so-called sigmoidal σ 's:

$$\sigma(t) \rightarrow \begin{cases} 1 & \text{as } t \rightarrow +\infty, \\ 0 & \text{as } t \rightarrow -\infty. \end{cases}$$

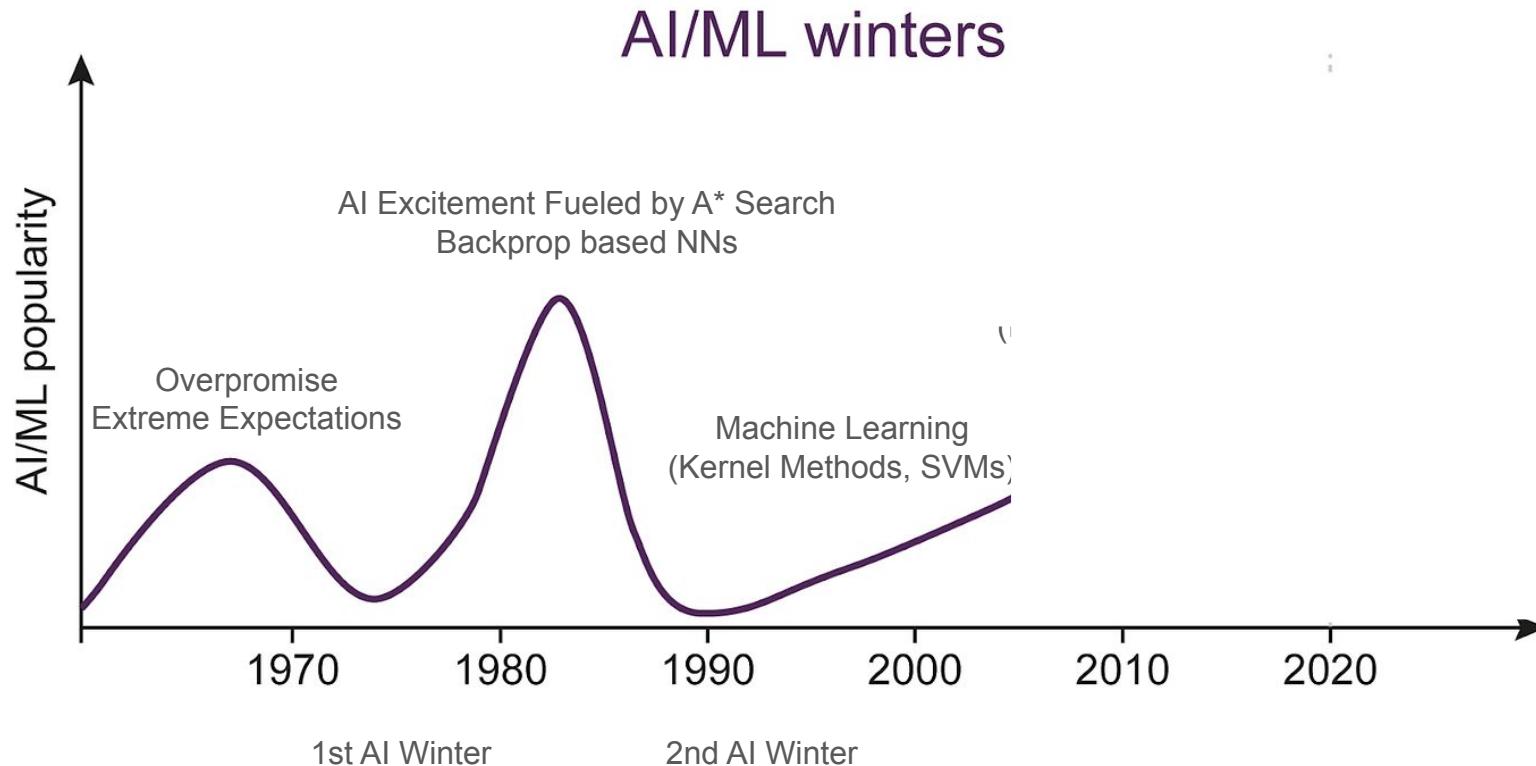
Such functions arise naturally in neural network theory as the activation function

Summer of SVMs 1995-2008

- 1993-1995 Corinna Cortes, Isabella Guyon, Vladimir Vapnik invent Support Vector Machines
- Mid 2000s ICML and NeuRIPS (NIPS) exclusively papers on non-neural network approaches
 - Mostly SVM, Graphical Models, Boosting
 - These algorithms are more efficient, easier to train / modify, have strong theoretical guarantees / frameworks



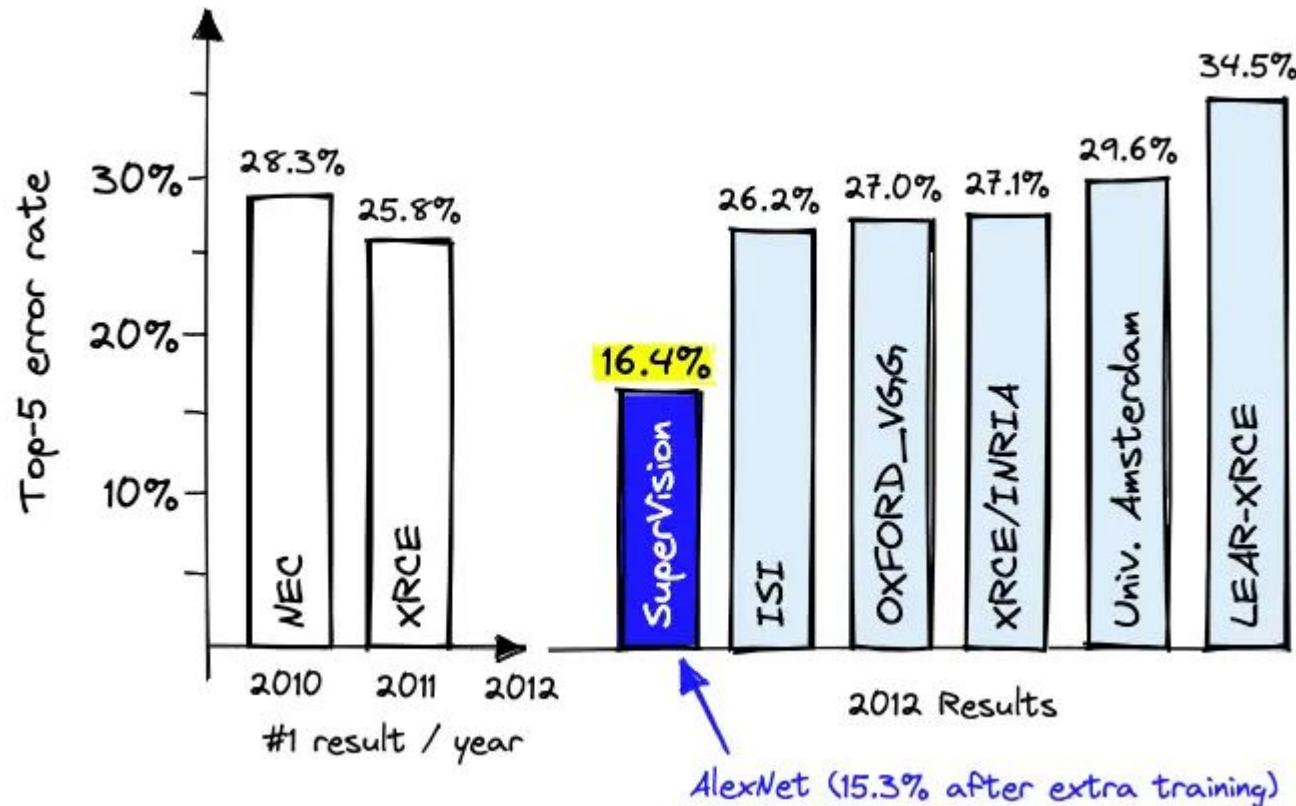
Public Perception of AI/ML



Neural Network Resurgence (2010s)

- Relentless effort by Hinton, Bengio, LeCun: Kept pushing Neural Nets when they were not cool - but did not join other communities (e.g. ICANN)
- Invent Deep Belief Nets in effort to attract experts in Graphical Models (mimics Graphical Models)
- Rename Neural Nets as “Deep Learning” (in effort to brand SVMs as “shallow”)
- Create ICLR as a venue to accept research on Neural Nets
- 2007 NeuRIPS Workshop on Deep Learning (rejected, changed to Hinton’s 60th birthday party)
- 2009 Fei-Fei Li creates ImageNet (after Caltech 4, 101, 256)
- 2012 Hinton’s deep network research creates AlexNet

Cornell Bowers CIS



Turing Award 2018



Yoshua Bengio



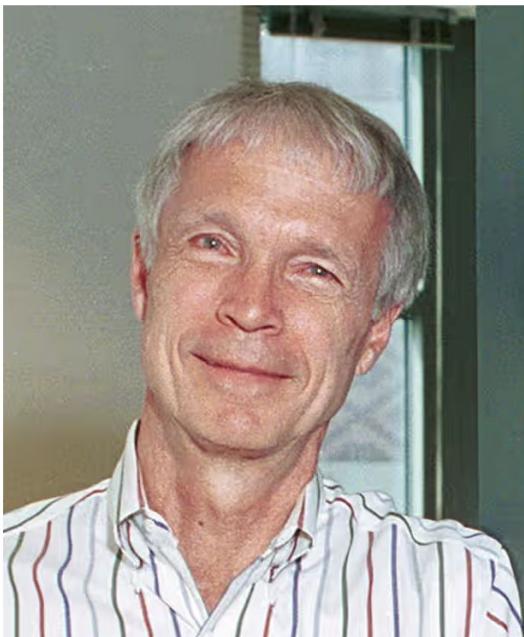
Geoffrey Hinton



Yann LeCun

Machine learning pioneers win Nobel prize in physics

Geoffrey Hinton, 'godfather of AI', and John Hopfield honoured for work on artificial neural networks



John Hopfield, left, and Geoffrey Hinton will share the 11m Swedish kronor (about £810,000) prize. Photograph: AP

Controversy



THE NOBEL PRIZE IN PHYSICS 2024 DECRIMINALIZES PLAGIARISM AND MISATTRIBUTION

IMAGE GENERATED WITH THE HELP OF A.I.

Jürgen Schmidhuber

Pronounce: You_again Shmidhoobuh
Technical Report IDSIA-24-24, IDSIA

AI Blog

Twitter: @SchmidhuberAI
7 Dec 2024

A Nobel Prize for Plagiarism



Jürgen Schmidhuber (2021, updated 2024)
Pronounce: You_again Shmidhoobuh

AI Blog
Twitter: @SchmidhuberAI

The most cited neural networks all build on work done in my labs

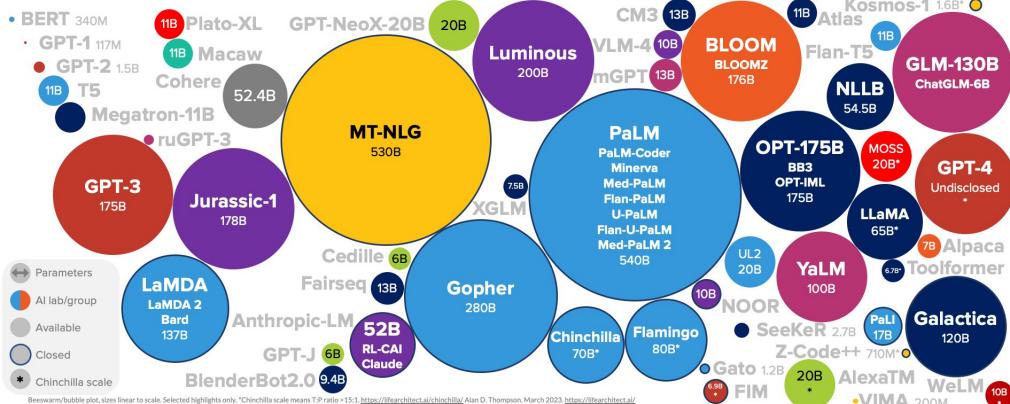
Abstract. Modern Artificial Intelligence is dominated by artificial neural networks (NNs) and deep learning [DL1-4]. Foundations of the most popular NNs originated in my labs at TU Munich



The Era of Scale (2020-Present)

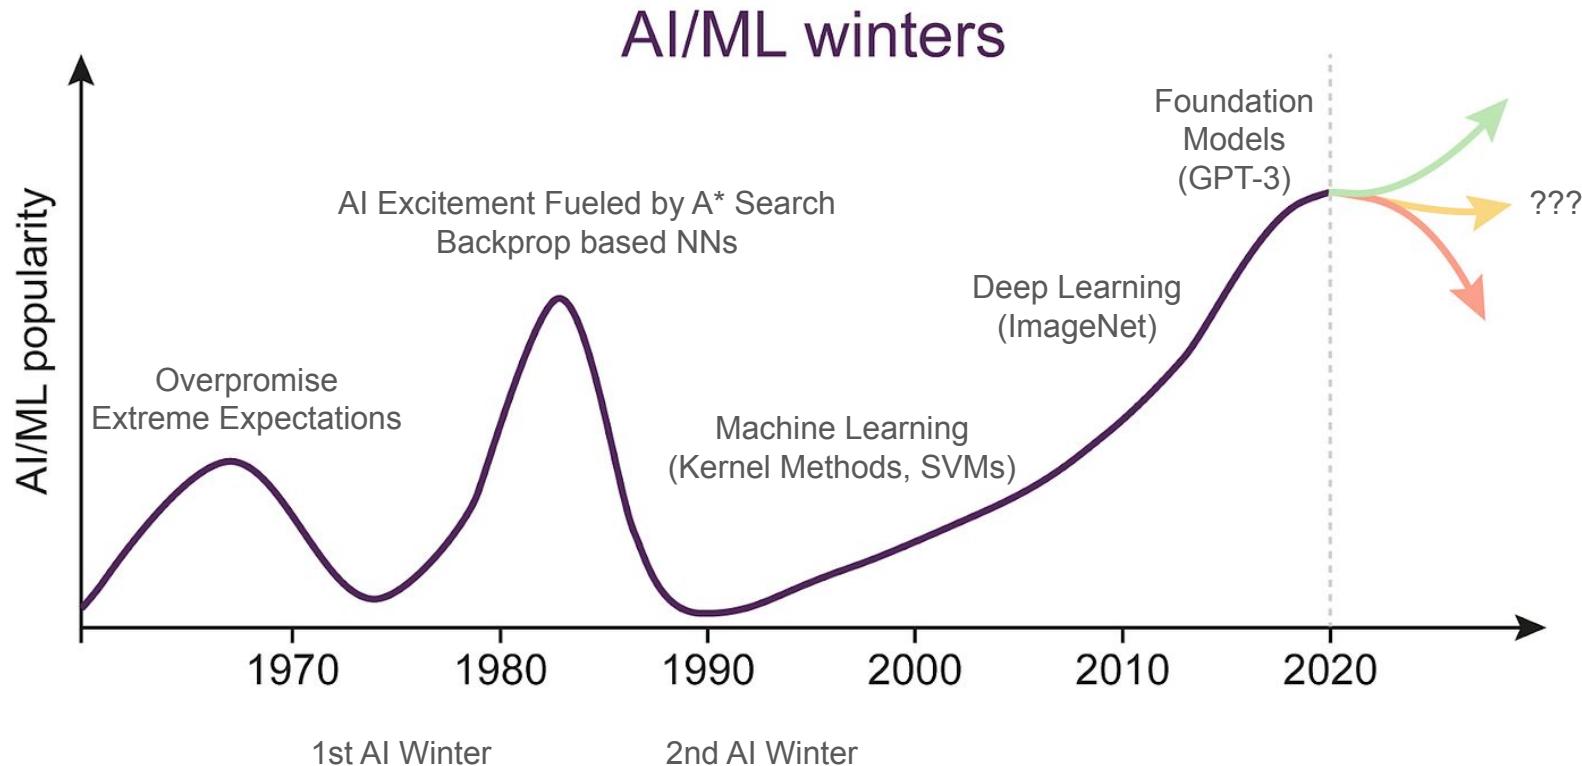
- GPT-3 introduced in 2020
 - “Language Models are Few-Shot Learners”
- Stable Diffusion released in 2022

LANGUAGE MODEL SIZES TO MAR/2023



<https://stability.ai/stable-image>

Public Perception of AI/ML



Task: Predict whether an image contains an eye.

Thanks!

- If you have received a permission number
 - Enroll today if you'd like to take the course
- Todo items:
 - Sign up for Google Colab Pro account (it is free for students)
 - Find study / project partners